

AI Output Verification Checklist

Trust but verify — every time

Claude is remarkably capable, but it's not infallible. It can hallucinate facts, invent citations, and present wrong answers with absolute confidence. This checklist is your quality gate — run every AI output through it before you use it.

Check 1: The Facts Check

Is this actually true?

DO THIS:

- Verify any specific claim — dates, statistics, names, events
- Google the key facts. If you can't find them, Claude may have invented them
- Check numbers for reasonableness. Does that revenue figure make sense? Is that date possible?
- Watch for confident-sounding nonsense — Claude never says "I'm not sure" when making things up

RED FLAGS

- Very specific statistics without a clear source
- Historical events you've never heard of
- Expert quotes that sound too perfect
- Technical claims that are almost but not quite right

Check 2: The Sources Check

Do these references actually exist?

DO THIS:

- Click every link. Dead links often mean invented sources
- Verify that cited authors actually wrote what's attributed to them
- Check that cited studies exist and actually say what Claude claims
- Look for "plausible-sounding" journal names that don't exist

RED FLAGS

- URLs that return 404
- Book titles that don't appear in any library catalog
- Papers attributed to real researchers that they never wrote
- Statistics cited to organizations that don't publish that data

Check 3: The Logic Check

Does this reasoning hold up?

DO THIS:

- Follow the argument step by step. Does each point actually support the conclusion?
- Look for circular reasoning — using the conclusion as evidence for itself
- Check for false equivalences and oversimplifications
- Ask: would an expert in this field agree with this reasoning?

RED FLAGS

- "Therefore" connecting two unrelated ideas
- Analogies that sound clever but don't actually apply
- Sweeping generalizations from a single example
- Presenting correlation as causation

Check 4: The Completeness Check

What's missing?

DO THIS:

- Ask: what would someone who disagrees with this say?
- Check for important caveats or exceptions that were left out
- Look for one-sided presentations of genuinely debatable topics
- Verify that "comprehensive" lists are actually comprehensive

RED FLAGS

- A pros-and-cons list where the cons feel weak
- "There are three main approaches..." (there might be five)
- Missing obvious counterarguments
- Recommendations without trade-offs

Check 5: The Freshness Check

Is this current?

DO THIS:

- Claude's training data has a cutoff date — it doesn't know recent events
- Verify any claims about current laws, regulations, or policies
- Check that recommended tools, websites, and services still exist
- Confirm pricing, availability, and feature claims independently

RED FLAGS

- "As of 2024..." or any time-stamped claim
- References to products or features that may have changed
- Legal or regulatory advice (always verify with current sources)
- Claims about who holds a position or role

The One-Minute Version

Short on time? Run these five questions before you use any AI output:

1	Can I verify the key facts with a quick search? If not, don't use them.
2	Do the sources actually exist? Click the links. Search the titles.
3	Does the logic hold if I push on it? Follow the reasoning, not the confidence.
4	What's being left out? Ask Claude directly: "What are the counterarguments?"
5	Could this be outdated? Check anything time-sensitive independently.

The Golden Rule: The more confident Claude sounds, the more carefully you should check. AI doesn't hedge when it's wrong — it states falsehoods with the same conviction as facts.

This checklist is a companion to Chapter 3 of [Creating with Claude: A Practical Guide for Everyone](#). In the book, you deliberately break Claude, catch a hallucination in the act, and build your own verification system. Get the book at marquepublishing.com.