

The Machine: Quick Reference Guides

Every file in your project, and a guide for every chapter's build

The book's appendix, for readers of Writing with Claude who signed up at marquepublishing.com. Part One is the artifact reference, one card per file. Part Two is the chapter-by-chapter process guide, one card per chapter. Print them, mark them up, give them to Claude.

Copyright © 2026 Chris Thatcher. Published by Marque Publishing, LLC. Free for readers of Writing with Claude. Not for resale or redistribution.



Part One — The Artifacts

These are the files that make up your machine. Each one exists for a reason, does a specific job, and has a specific format. They live in your project folder and travel with you from book to book.

Before the individual cards, here is how the files talk to each other. This circulation is the machine:

Every time something goes wrong in a session, you capture it in the **scar journal**. The pattern becomes a rule in the **rules file**. The rules file loads in **preflight** alongside **Canon**, your **voice profile**, the last entry from your **session log**, and the bar. Claude reports back, you say go, and you draft. You edit against your **benchmark passage**. You close the session with **postflight** — the six craft questions, the log entry, the **progress tracker** update. Tomorrow, you open by reading that log entry, and the loop starts again. Every session, the files get better. Every session, the AI gets closer to what you need.

That's it. That's the whole system. The rest of this section tells you what goes in each file.

Canon

Canon is the source of truth for your book — every locked decision, in one place, loaded every session.

Without it, decisions live in conversation history that the AI can't see tomorrow. You and Claude spend an hour working through whether your protagonist's mother is alive or dead, you commit to dead, and next session the AI has no idea that conversation happened. Canon is the file that makes decisions persist.

What goes in it:

- Your thesis — the one sentence your book exists to deliver. First entry, always.
- Your north star authors — the three writers whose company your book deserves to keep.
- Every structural decision from scaffolding — chapter order, arc, character assignments.
- Every locked decision made during drafting — plot points, arguments, character details, anything you've committed to.

How to build it: Tell Claude to create a file called Canon in your project folder. Your thesis is the first entry. From that point forward, every time you make a decision you won't reverse, tell Claude to write it into Canon.

How to use it: Load it every session as part of preflight. When you lock a new decision during a session, add it before you close.

The prose-is-primary rule: When Canon and the prose conflict, the prose wins. Canon records intent, but the shipped page is the final word. This rule exists because early in my second book, a backstory document overwrote a decision the prose had already made. The prose is the primary canon. Always.

Canon is not a notebook — it's a registry. Only locked decisions go in. If you're still deciding, it stays out until you commit. It grows slowly and deliberately, and it travels with you to the next project.

Session Log

The session log is your memory between sessions. Every working session gets an entry. Every entry follows the same shape.

The AI has no memory of prior sessions unless you load one. You could have the best session of the project — solve a structural problem, lock three decisions, draft two thousand words — and tomorrow the AI will greet you like a stranger. The session log is what bridges the gap. The last entry tells the next session what happened, what's open, and where to start.

What goes in each entry:

- Date
- What happened — work completed this session
- What was decided — decisions locked, choices committed
- What is open — unresolved questions, pending tasks
- What comes next — the specific target for the next session

Once you're drafting (Chapter 9 forward), add the six craft questions: Did it land? Are we at the bar? Would our north star authors sign this? What did we get right? What could we have done better? What did we miss?

How to build it: Have Claude create a file called Session-Log in your project folder. Write your first entry tonight.

The one rule: Every session opens by loading the last entry — not the whole log, just the last entry. Every session closes by writing a new one. The log is the handoff between tonight and tomorrow. Skip it once and the next session starts cold.

Plain text, chronological. Date header, then prose. No template required — the consistency comes from writing one every session, not from formatting.

Scar Journal

The scar journal is the evidence file. Every time the AI does something wrong — invents a detail, oversteps scope, loses work, gives a confidently wrong answer — you write it here.

Rules without evidence are arbitrary. The scar journal is why each rule exists. When you read it front to back, the pattern emerges — the same kinds of failures, recurring. That pattern is what builds your rules file.

What goes in it: What happened — the specific failure. When it happened — which session, which chapter. What you caught — the detail that was wrong, the scope that was overstepped, the work that vanished. And most importantly, what you almost let through. The near-miss is often more important than the catch, because the near-miss is the one that would have shipped.

How to build it: Claude creates it on your first night (Chapter 1), before anything goes wrong. You're building the file before you need it, so when you do need it, it's already there. When something feels off in a session — a wrong fact, a strange answer, something you didn't ask for — stop and have Claude capture the context immediately. Don't wait until the end of the session. The detail matters, and detail fades.

The pipeline: The scar journal is the evidence. The rules file is the verdict. A scar becomes a rule when you can write one sentence that would have prevented it. The journal feeds the rules file. The rules file feeds the preflight. That's how tonight's failure makes next week's session better.

Rules File

The rules file is a text file with plain-language rules that Claude loads at the start of every session. It's the accumulated wisdom of everything that has gone wrong, written as instructions to prevent it from going wrong again.

The AI will make the same mistake twice unless you tell it not to. Tell it once in conversation and the fix lasts one session. Write it in the rules file and the fix lasts forever. This file grows for the rest of your writing life.

What goes in it:

Scar-born rules — one sentence per failure pattern. Not code. Not configuration. Sentences, written in the language you would use to tell a person what to do. Example: *"Do not invent details that are not in the source material. Tag anything questionable with [TO CONFIRM]."*

Standing orders for preflight: *"At the start of every session, before drafting, load: rules file, Canon, voice profile, last session log entry. Confirm the target. State the bar. Report back what you loaded and wait for my go."*

Standing orders for postflight: *"At the close of every session: log the session, save all artifacts to the project folder, update Canon with new decisions, update the progress tracker."*

The scope rule: Do not rewrite content you were not asked to rewrite. Do not change meaning or voice.

The evidence rule (from Chapter 11): *"Nothing about you, or as you, ships unverified."*

The banned-phrase list — AI tics caught and killed during editing. This list lives inside the rules file and grows with every edit pass. Seed it from Chapter 7's fingerprint patterns and expand it every time you run Pass 1 in Chapter 10.

How to build it: Create it in Chapter 6. Seed it with the standing orders and any rules your scar journal has already earned.

How to use it: Load it first, every session, as step 1 of preflight. The full file, not a summary, not a search. Every rule, every session, no exceptions.

New scar, new rule. The file gets longer. The sessions get better.

Voice Profile

Your voice profile is your voice described in your own words — loaded every session so Claude knows how you write instead of guessing.

Without it, Claude writes in its own voice — competent, clean, and generic. The kind of prose that passes a checklist but doesn't pass a reader. With the profile loaded, Claude gets close enough to edit, close enough to make yours. That difference is the whole point.

What goes in it:

Rhythm — your sentence-length patterns. Short, long, mixed. Where do you break? How does your pacing work? My sentences are uneven on purpose. I run a thought for forty words when the argument needs room, and then land it in five. That's mine. Yours is different. Name it.

Diction — formal or conversational? Technical or plain? What words and phrases are distinctly yours? What do you reach for when the writing is working?

Refusals — what you never say, what you never do. The banned-phrase list starts here. The negative space is as much the voice as the words. I will never write "unlock value" or "transformation journey." Those are mine. Name yours.

If you write in more than one register — professional versus personal, fiction versus nonfiction — name each one. Write what changes between registers and what stays constant. The constants are your voice. The variables are your registers.

Who writes it: Claude drafts it from your samples. You rewrite it. Claude's draft will be accurate and a little clinical — a linguistics paper about a person instead of a person describing how they talk. Rewrite it the way you would describe your writing to a friend. Five sentences, maybe ten, to start. It grows from there. Two things you never do: let Claude guess at your voice without samples, or accept its draft unedited.

How to build it: Gather at least three writing samples from before you used AI — the pieces that sound most like you on a good day. Upload them. Have Claude analyze the patterns and draft a profile from what it finds. Then rewrite the profile in your own words.

How to calibrate it: The calibration that surprised me: I naturally use some of the same patterns the AI overuses. I write parallel constructions. I use anaphoric runs. I land on short punchy closers. The difference is frequency and placement — I do it once a chapter where the AI does it every section. The principle: ban the saturation, not the existence. Calibrate for your natural rate.

How to test it: Take a piece Claude drafted before the profile existed. Rewrite it with the profile loaded. If the difference isn't obvious, the profile needs more specificity. Sharpen it and test again.

Benchmark Passage

The benchmark is a passage you wrote yourself — in your voice, from scratch, without AI assistance. It's the fixed bar every draft gets measured against.

The voice profile tells Claude how you write. The benchmark proves it. When you're editing and you need to know whether a paragraph sounds like you or sounds like anyone, hold it next to the benchmark. The answer is immediate.

Write an opening passage — the first scene, the first argument, the first story — rough and fast, in your voice. Then have Claude draft the same passage cold, with no exposure to yours. Put them side by side. Yours wins. Name the specific quality that makes it yours — your compression, your rhythm, your way of entering a scene. That quality is the pattern that is yours and only yours. Give it to Claude. Tell Claude: *"This is the bar. Every draft gets measured against this passage."*

The benchmark is fixed. You don't rewrite it as your skills improve — you wrote it at the start as proof of what your voice sounds like when nobody is helping. It stays where it is, doing its job, for the rest of the project. In Chapter 10's Pass 3 (voice calibration), you read every paragraph against the benchmark and mark the ones that don't match the standard.

Facts File

The facts file is everything true about you that belongs in an author bio or public-facing material, sourced from evidence.

The AI fabricates author details with confidence. The bio test in Chapter 11 proves it — open a clean conversation, ask Claude to write your bio, and count the inventions. There will be inventions. Credentials you don't have. Books you didn't write. Positions you never held. All delivered with the same confidence as the facts it gets right.

The facts file is the cure. Everything true about you that belongs in a bio — career, publications, credentials, location, personal details you choose to share. The evidence for each fact — your resume, your LinkedIn, your publication history. Nothing unverified. Nothing assumed. Nothing Claude told you about yourself.

Build it once. Write your About the Author page from it. Copy your bio to every platform from it. Check every piece of public-facing text with your name near it against it. One file, verified once, used for the rest of your writing life. Add the rule to your rules file: *"Nothing about you, or as you, ships unverified."*

The Supporting Files

Three more files round out your project folder. They don't need the full treatment — they're simpler tools that do one job each.

Progress tracker. A table — chapter, title, status — updated every session as part of postflight. Create it when you create the rules file in Chapter 6. When you come back tomorrow and ask "where do I stand?" this is the answer. The format doesn't matter. The habit of updating it does.

Interior spec. One page, eight decisions that control your entire interior: trim size, typeface, margins, running heads, page numbers, chapter openings, matter order, widow/orphan control. Lock it before typesetting begins — without a locked spec, every interior decision gets made ad hoc, and the result looks like it. Fiction and nonfiction are different systems — fiction is first-line indent and justified; nonfiction is block paragraphs and left-aligned. Use the Interior Spec Worksheet from the production templates for the full template with geometry checks and a trap list.

Production worksheets. The owner-column checklist and data sheets that run your production — ISBNs, metadata, copyright page, cover wrap, platform uploads, deposit and registration. Production has two lanes: Claude's (pre-fill data, compute numbers, draft documents) and yours (account actions that require your logins). The worksheets keep the lanes clear. Download the template set from the Marque Publishing website or use the files in the `production_templates` folder. Give the Production Run Sheet to Claude and tell it to run your production. The sitting pattern is simple: data sheet on the left, platform on the right, work the fields in.

The full set: Run Sheet, ISBN Log, Interior Spec Worksheet, Copyright Page Template, Metadata Worksheet, Cover Wrap Spec, Platform Upload Checklist, Deposit and Registration sheet, Editorial Brief, Cover Brief.

The Briefs

Two documents you write before you hire anyone — and keep even if the anyone turns out to be Claude. The editorial brief goes to editors before the edit (Chapter 10). The cover brief goes to the designer before a single image exists (Chapter 11). Each one says what the job is, what it is not, what done looks like, and the one rule that disqualifies a candidate. That last line is what makes an audition scoreable instead of a matter of taste. Both are in the download set as fill-in templates; here is the shape.

The editorial brief. A cover note in your own voice, including the fact that AI was used in drafting and what has already been done about it. An at-a-glance table: title, genre, length, audience, format, delivery and ship dates. About the author, from the facts file. The voice section — the part that matters most — with two lists drawn from your voice profile: the tics that stay (the things that look like mistakes and are not) and the moves to push back on (where your prose drifts when it drifts). Scope: what you are asking for, what you are not, and the disqualifier in the exact words you will score against — *an edit is not a rewrite*. Style conventions as a table. Timeline. Deliverables: tracked-changes manuscript, explicit queries, a style sheet, an optional one-page memo. How to respond. One representative sample chapter at the end.

The cover brief. Project at a glance with title and subtitle marked as locked text. Audience, with three books they already own. What the book actually is, in your words, plus the two or three signature ideas a concept can anchor on and how literally each may be shown. Tone, with a comparison a designer can picture. The two jobs the cover has to do at once — credibility in the hand and presence at thumbnail size — and how you want them balanced. Reference covers, split by what to learn from each. The avoid list: the cliches of your shelf, named specifically. This is the disqualifier; a concept that uses one ends the conversation. The authenticity rule: real things verified, or nothing. Deliverables (ebook front, print wrap, layered source files). Technical specs from the cover wrap worksheet, marked provisional until the page count freezes. Rounds, and the exit criteria for any modify-the-existing request. How to respond.

How to build them: Have Claude draft both from your Canon, voice profile, facts file, and rules file. Then rewrite the parts that are you talking — the cover note and voice section of the editorial brief, the what-the-book-is and tone sections of the cover brief. If those sound like anyone else's book, you will get anyone else's edit and anyone else's cover.

Project Folder Structure

By the time you're drafting, your project folder will have dozens of files — chapter drafts in multiple versions, quarry files, source captures, project infrastructure, and production artifacts. Keep them flat and you'll spend more time finding files than writing. Organize them early.

The structure:

Folder	What goes in it
drafts/	Every chapter draft — all versions, .md and .docx. Introduction, conclusion, and appendix drafts too. This is where the prose lives.
quarries/	Quarry files — the raw material mined from session logs and conversations that each chapter drafts from.
sources/	Source captures — conversation transcripts, research notes, reference documents you brought in from outside the project.
reference/	North star samples, archived project files, cross-chapter audits, the Reader-Thread-Map, anything you consult but don't edit.
workshop/	Workshop materials, if you're building one alongside the book.
cover/	Cover files — concepts, working files, final artwork.
production_templates/	The production worksheet set from Chapter 12, plus the editorial and cover brief templates.
_archive/	Superseded versions, abandoned approaches, anything you're done with but want to keep.

What stays at root: The infrastructure files that preflight loads and postflight updates — Project-Instructions, Drafting_Rules, Canon, Rules File, Session Log, Scar Journal, Voice Profile, Benchmark, Progress Tracker, Facts File, Interior Spec, Blueprint. These are the files you and Claude touch every session. They stay where you can see them.

When to set it up: Chapter 1 or Chapter 2 — before the file count makes it painful. Tell Claude to create the folder structure and move files into it. The earlier you organize, the less archaeology you do later.

Part Two — The Chapter Guides

Each guide matches a chapter. On your first book, work them in order. On your second book, skip to whichever guide you need.

Chapter 1 — Getting Started

What you build: Your project folder and your first brain dump into Claude.

Set up the folder:

Create a folder on your computer for this book project. Name it something you'll recognize in six months. Dump everything into it — notes, emails you sent yourself, half-written chapters, fragments, voice memos transcribed, research links, photos of napkin sketches. Do not organize. Do not clean up. Get it all in one place.

Set up the project:

Open Claude. Create a new project and point it at your folder. Start talking. Use the microphone if that's easier. Talk the way you would talk to a friend who finally asked about the book you've always wanted to write. Cover all of it — the idea, the people in it, the thing you want to say, the parts you've already written, the parts still stuck in your head.

Be specific. Be verbose. Fragments are fine. Contradictions are fine. Get every piece out of your head and into the project.

When you run out of things to say, tell Claude: *"Summarize everything I just told you about this book. Organize it by topic. Tell me what's clear and what's still vague."*

Read the summary. Correct what's wrong. Add what's missing. Push until the summary matches what's in your head.

Create the session log:

Create a file called `Session-Log` in your project folder. Log tonight's entry — date, what happened, what's still open, what to do next session. This file is your memory between sessions. Save it before you close. Every session from here forward opens by reading the last entry and closes by writing a new one.

What you should have when you're done: A project folder with your raw material. A Claude project pointed at that folder. A first conversation where you got everything out and Claude organized it back. A session log with tonight's entry. This is your quarry — the raw material every chapter mines from here.

Chapter 2 — Your Thesis

What you build: The one sentence your book exists to deliver. The Canon file.

Build the thesis:

Open your project. Tell Claude you're ready to build your thesis. Answer the interrogation questions. Write your answers down — don't just think them:

- **Who is this book for?** Name the specific person, not a demographic.
- **What do they believe right now that you're going to change?** The wrong belief your book corrects.
- **What is the one sentence on the back cover that makes someone pick this up?** Not a topic. A promise.
- **Where does the reader start and where do they finish?** The transformation — what they walk in believing and what they walk out knowing.

Your first answers will be vague. Do not accept them. Keep refining. An idea is "I want to write a book about X." A thesis is a specific promise to a specific reader with a specific transformation. When one sentence survives everything you and Claude can throw at it — and you would put it on the back cover tonight — that's your thesis.

Create the Canon file:

Create a file called canon in your project folder. Your thesis is the first entry. Pick three authors whose company your book deserves to keep. These are your north star authors. Write their names in Canon. Every session ends by asking: would these authors put their name on this page? Every locked decision from here forward goes into Canon. It is the source of truth for your book.

Update the session log:

Open the session log you created in Chapter 1. Log tonight's entry — your thesis, any decisions you locked, what to do next session. The log has been running since your first session and it keeps running from here.

What you should have when you're done: A thesis that survived interrogation. A Canon file with the thesis and your north star authors. A session log updated with tonight's entry.

Chapter 3 — Your Scaffold

What you build: The chapter-by-chapter structure of your book, with every decision locked.

Open the session:

Open your project. Read the last session log entry. Ask Claude where you left off. Then give Claude the scaffolding brief — Claude does not know what you mean by scaffolding until you tell it:

"We are going to build the scaffolding for my book before any drafting. Scaffolding means: the thesis (one sentence — a specific promise to a specific reader), the arc (where the book starts, where it ends, and what changes in between), a chapter map where every chapter has a job to do for the thesis, and for fiction a Goal, Motivation, Conflict profile for every major character. I make every decision. You ask questions, offer options, and challenge my thinking. Every decision I lock gets written to Canon. Do not draft prose. Start by asking me about the arc, and hold my thesis up as the standard for everything we build."

Claude proposes. You accept, reject, or redirect. You set the direction. You and Claude go back and forth developing the story until you are satisfied, then you ask Claude the six questions: did it land, are we at the bar, would our north star authors be proud of this, what did we get right, what could we have done better, and what did we miss.

If there was one thing I can point at across four finished books and two more in the works, that made my work consistently better and more of my own, this set of questions is it. And you take the time to go back through every single item that Claude flagged. Don't be satisfied until you are at the bar. Treat every flag as an opportunity to make the work better. This is how you get the end of the scaffolding build and feel like you already wrote your book.

If your book is fiction:

Start with your people. Build a Goal, Motivation, Conflict profile for every major character — **Goal:** what do they want, **Motivation:** why do they want it, **Conflict:** what stands in the way. Without GMC you don't have a story, you have a situation. Use a placeholder for anything unresolved — fill it in later, but don't skip the slot.

Build your story arc — the events in order, the turning points, the resolution. Build your chapter map — which scenes live in which chapters, what each chapter does for the story.

If your book is nonfiction:

Start with your stories — the real ones only you can tell. Run story-mining sessions: tell Claude the experiences, the failures, the moments that taught you the thing your book teaches. Then the thing you want your reader to take away — chapter by chapter, the path from where they start (your thesis) to where they finish. Build your chapter map. Each chapter advances the reader one step toward the promise in your thesis. If a chapter doesn't advance them, it doesn't belong.

Lock decisions:

When the chapter map exists, walk it one chapter at a time. Tell Claude you are ready to walk Chapter 1. Claude proposes the beats — for fiction, what happens in the scene and who wants what; for nonfiction, the story that opens the chapter, the lesson it carries, and what the reader does with it. You accept, reject, or reshape each beat. When a beat is right, say lock it and Claude writes it into Canon. Before you leave the chapter, have Claude cross-check what you locked against everything already in Canon — names, dates, timeline, decisions — and flag conflicts. Then walk the next chapter. Log every session. You are done when the book feels already written.

What you should have when you're done: A chapter-by-chapter scaffold with the arc mapped. Every structural decision locked in Canon. Session log current.

Chapter 4 — The Reader-Lens Walk

What you build: Proof that your scaffold works for the reader, not just for you.

The walk:

Open your project. Have Claude read the last session log. Before you move anything, tell Claude to save a dated backup of your scaffold. Name it with today's date. This is your insurance policy — you can always get back to the pre-walk version.

Walk the scaffold from Chapter 1 to the end, but as your intended reader — not as the author. You know what every piece means, so every order makes sense to you. Your reader doesn't. At each chapter, ask: **Does my reader have what they need to understand this?** If the answer requires something that comes later, the order is wrong. The question that matters: **Am I taking my intended reader on the right journey?**

When something is out of order:

Move the pieces where your reader needs them. Fix every pointer that now aims at the wrong door — callbacks to chapters that moved, setups that now pay off in a different place. Rewrite the connective tissue so each chapter hands off clean to the next. Lock every change into Canon with the reasoning behind it. Log the session.

What you should have when you're done: A scaffold that works for your reader. A dated backup of the pre-walk version. Canon updated with every change and why you made it.

Chapter 5 — Zero Trust

What you build: The understanding that the AI's memory is in your files, not in the AI. The scar journal.

The amnesia test:

Open a fresh chat window — not your project. A clean conversation with no context loaded. Type: *"What do you know about my book?"* The answer is nothing. The collaborator from four chapters does not exist in that window.

Now go back to your project. Load your rules file, Canon, and the last session log entry. Ask the same question. Watch the difference. The memory is in the files. It was always in the files.

The scar journal:

Create a file called `Scar-Journal` in your project folder if you haven't already. Read it from the first entry to the last. The pattern will emerge — the same kinds of failures recurring. This is important, because it makes your system better every time you add to it. When something goes sideways, ask Claude to capture the context in the scar journal and ask what rule would have prevented it. The rules go into the next file. That won't guarantee you don't get bit by the same thing later, but you give Claude clear instructions to catch the problems before they surface. These are the named failure modes to watch for:

- **Fabrication** — invented details that sound plausible and survive multiple drafts because you don't catch them
- **Canon violation** — reaching back to superseded versions instead of the locked current version
- **Partial-context reads** — acting on a grep hit without reading the paragraph around it
- **Amnesia** — the model has no memory of prior sessions unless you load it
- **False denial** — inventing technical excuses for why something can't be done

Every time you catch a new failure, write it in the scar journal: what happened, when, what you almost let through. The next chapter turns each scar into a rule.

What you should have when you're done: The conviction that trust belongs in the files, not in the AI. A scar journal with your failure patterns visible. The foundation for everything in Chapter 6.

Chapter 6 — The Trust Stack

What you build: The enforcement loop that runs every session for the rest of your writing life.

Create the rules file:

Create a file in your project folder. This is your rules file — plain-language rules, one per line. Not code. Sentences Claude can read and follow. Seed it with the standing orders below and any rules your scar journal has already earned.

The preflight (run this at the start of every session):

Tell Claude to run these five steps, in this order. First, **load the rules file** — the full file, not a summary, not a grep. Every rule, every session. Second, **load Canon** — every locked decision. If Canon and the prose ever conflict, the prose is primary. Third, **load the last session log entry** — not the whole log. Just the last entry, which tells you what happened, what's open, and what to do next. Fourth, **confirm the target** — state aloud what you're going to draft or revise this session. One chapter, one task, stated clearly. Fifth, **report back and wait** — Claude gives you a short brief of what it loaded, where you left off, what is open, and today's target, then waits for your go. Nothing gets drafted until you say so. This is where you catch the stale file or the missed requirement before it costs you a chapter.

After Chapter 7, add your voice profile to the preflight load. After Chapter 8, add the bar — Claude states your benchmark, your north star authors, and whatever the last postflight said was missing before the drafting starts.

The postflight (run this at the end of every session):

Log the session — what happened, what you decided, what's still open, what to do next session. **Save every artifact** to the project folder. Work that is not saved is work that did not happen. **Update Canon** if you locked any new decisions. **Update your progress tracker**.

The scar-to-rule pipeline:

When Claude invents a detail, oversteps scope, or does something you didn't ask for: write the scar in the journal — what happened, when, what you caught. Write the rule in the rules file — a plain-language sentence. Example: *"Do not invent details that are not in the source material. Tag anything questionable with [TO CONFIRM]."* The next session is better because the rules file is longer.

New scar, new rule. No exceptions.

Make preflight and postflight a skill. If you're using Claude Code, Claude Desktop, or any tool that supports skills, save preflight and postflight as named skills — reusable instructions Claude loads on command. Preflight as a skill means you say "run preflight" and Claude loads the rules file, Canon, voice profile, and last session log entry, confirms the target, and reports back without you listing the steps. Postflight as a skill means you say "run postflight" and Claude runs the six craft questions, logs the session, updates Canon and the progress tracker. The skill turns a checklist into a single command. When you add a step to either one, tell Claude to add it to the skill. Build them once, use them every session, every project.

What you should have when you're done: A rules file with your scar-born rules and the preflight/postflight standing orders. Preflight and postflight running as named skills. The loop running every session from tonight forward.

Chapter 7 — Your Voice Profile

What you build: Your voice, on file, loaded every session as part of preflight.

Identify the AI fingerprints (what to strip):

Before you define your voice, know what AI voice sounds like so you can ban the saturation, not the existence:

- Perfect parallel triples where every clause carries the same weight
- Aphoristic three-word section closers that sound profound and say nothing
- Anaphoric runs — three or four sentences starting with the same word
- Comma-stacked noun lists of four to six items
- Tidy-clever phrases that read as model cleverness instead of human insight
- Reframe-via-negation pairs ("This isn't about X. It's about Y.") used as structural crutches

Your natural writing may include some of these occasionally. The AI overuses them. Ban the saturation rate, not the device itself.

Build the profile:

Gather at least three writing samples that predate your AI use. The pieces that sound most like you on a good day. Upload them to Claude. Ask Claude to analyze the patterns — sentence length, rhythm, diction, what you do that's distinctive.

Have Claude draft the profile from what it found. Then **rewrite it yourself**, the way you would describe your writing to a friend. Claude's draft is the accurate half; your rewrite is the half that sounds like a person. Cover at minimum:

- **Rhythm** — short sentences, long sentences, mixed? Where do you break lines? How does your pacing work?
- **Diction** — formal or conversational? Technical or plain? What words are yours? What phrases do you reach for?
- **Refusals** — what you never say, what you never do. The banned-phrase list starts here and grows in Chapter 10.
- **What you sound like when the writing is working** — describe it in your own words.

If you write in more than one register (professional vs. personal, fiction vs. nonfiction), name them. Write what changes between registers and what stays constant. Tell Claude which register applies to this project. Save the profile to the project folder. Add the voice profile to your preflight load: rules, Canon, voice, last session log, confirm the target, report back.

Test it:

Run preflight and confirm the voice profile shows up in Claude's brief. Then take a piece Claude drafted during the early sessions — before the voice profile existed. Ask Claude to rewrite it with the voice profile loaded. Read the rewrite against the original. If the difference is not obvious, the profile needs more specificity. Sharpen it and test again. Then run it the other direction: hand Claude a paragraph you wrote and ask what in it does not match the profile. If it flags something that is you, the profile is missing a line. Add it.

What you should have when you're done: A voice profile in your own words. The profile loaded as part of preflight. A test that proved it works.

Chapter 8 — Your Benchmark

What you build: Your own writing as the fixed bar every draft gets measured against.

Write it yourself:

Pick the chapter or section you've been scaffolding. Open it. Write the opening — the first scene, the first argument, the first story. Write it rough and fast in your voice. Do not ask Claude for a draft. Do not let Claude touch it yet.

Run the side-by-side test:

Hand Claude the same assignment — cold. No context beyond the project files, and no exposure to your version. Ask Claude to draft the same passage from scratch. Put them side by side. Read both against your voice profile. Yours will win.

Name the pattern:

Identify the specific quality that makes your version sound like you wrote it. It's not "better" — it's something specific. Your compression, your rhythm, your way of entering a scene. Write it down. Give it to Claude. Tell Claude this is the thing that must survive every edit.

Set the bar:

Hand Claude your version now. Let Claude tighten it — flag soft spots, question structure, cut redundancy. That is Claude's job. Your job was writing it. Claude's job is making it tighter without making it less yours. The passage that survives is your benchmark. Tell Claude: *"This is the bar. Every draft I write and every edit you suggest gets measured against this passage."* Then have Claude add the bar to the preflight skill. From now on, after it loads rules, Canon, voice, and the last log, it states the bar for the session — your benchmark, your north star authors, and whatever last night's postflight said was missing — before anything gets drafted. Everything Claude writes is measured against the bar before you ever read it.

What you should have when you're done: A benchmark passage in your own voice. The bar set in your own words. Proof that the pen is yours. The bar loaded in preflight for every drafting session from here forward.

Chapter 9 — The Drafting Cadence

What you build: The repeatable session that writes your book, one evening at a time.

Every session, the same shape:

Run Preflight every time you start a new chapter. Load the stack: rules, Canon, voice profile, last session log entry, the bar. Confirm the target — one chapter, one section, one task. Wait for the brief before anything gets drafted. Not "as much as I can get done." One target. If you don't finish the chapter in a single sitting, tell Claude when you are done for the night and tell it to write the session log to capture the work. When you come back, ask Claude where you left off and if you need to, ask it for a summary of the last session.

Draft in conversation. Two modes, and you'll use both. If the words are in you, write the passage yourself — rough, fast, in-voice. Then hand it to Claude and let Claude tighten it. Watch the word counts fall as it gets closer to you. If the words aren't coming, direct Claude with specificity. Not "write the next section" — tell Claude what the section does, what happens in it, what it should feel like. Then rewrite what comes back until it passes the benchmark.

Set one target per session. A scene, a section, an argument. Chapters accumulate. That's the trick — not a heroic month, a repeatable evening, run enough evenings in a row.

Run Postflight every time you finish what you think is a solid draft. Log what happened. Answer the six craft questions in writing:

- **Did it land?** Does the reader feel what this chapter was built to deliver?
- **Are we at the bar?** Would you read this aloud and be proud of it?
- **Would our north star authors sign this?** Load the names from Canon. Apply the test.
- **What did we get right?** Name the specific beats, lines, and choices that worked.
- **What could we have done better?** Voice slippage, sections that explain instead of show, anything generic.
- **What did we miss?** Missed callbacks, canon gaps, requirements from the last session log that didn't land.

Pay attention to the things Claude flags and don't be satisfied until you addressed the things you missed or could have done better. Tell Claude when you are satisfied and lock the chapter as green. Save everything.

When a chapter fights you: Look for why. Were you not feeling it and had Claude draft it for you, only to find it totally missed the mark? Sometimes it is easier to just tell Claude, this is crap and not even close to the quality of the work we have been doing. It will go back and assess the output and most likely say you are right. Pick out the things that need to be fixed and discuss the ideas with Claude until you are satisfied. Sometimes you will rewrite the whole thing and it will be better. Sometimes, you and Claude come up with new ideas and it drafts something that reflects the quality you have been creating. And sometimes, you just say, throw it out and let's start again.

What you should have when you're done: A finished draft, built one session at a time, with every session logged and every chapter measured against the bar.

Chapter 10 — The Five-Pass Edit Machine

Your draft is done, but there is lots of work left to do. Editing is next up on the list. This is one of the places you could choose a person instead of Claude. I told you my experience with editors changing my voice and my meaning. I chose Claude. I created a five-pass approach to editing because not every edit pass is focused on the same thing. Rather than give Claude too many things to focus on at once, I created a pass for each function.

What you build: The editing system that makes your draft publishable — in your voice.

If you are auditioning a person: write the editorial brief first (template 08_Editorial_Brief_Template) and score every sample against its disqualifier — an edit is not a rewrite. If Claude is the editor, the same brief is its instructions.

Run these five passes on every chapter, in order. Each pass has one job. Do not combine them. Not every pass needs the most expensive model; the assignments below are what I use, and they will change as the models do.

Pass 1 — Fingerprint scan. *Model: Sonnet.*

Have Claude scan the chapter against your rules file and every tic on your banned-phrase list. The scan is not a rewrite — Claude flags each hit with the paragraph before and after for context. Walk each flag interactively:

Read the flag in context. Rule it: keep (it's your voice, not a tic), cut, or rewrite. Push back when you disagree. The ruling is always yours. Every new tic you kill gets added to the banned-phrase list in your rules file. The list grows for the rest of your writing life.

Pass 2 — Continuity. *Model: Opus.*

This one has to run across the entire manuscript. Every name, date, place, age, object, and fact checked against Canon and every other chapter in the manuscript. Claude cross-references — but when Claude reports a conflict, you read both passages yourself. The prose is the primary canon.

Check specifically: character names spelled consistently, dates and timelines that add up, ages that don't contradict, objects that don't teleport, places described the same way twice.

Pass 3 — Voice calibration. *Model: Opus.*

Read the chapter against your benchmark passage from Chapter 8. Not for errors — for register. Mark every paragraph that doesn't sound like the standard. Fix them yourself or direct the fix, but you make the calls. This pass is yours.

Ask: "Does this paragraph sound like the benchmark, or does it sound like anyone could have written it?"

Pass 4 — Copy edit. *Model: Sonnet.*

Punctuation, grammar, spelling, formatting. Under the standing rule — add it to your rules file before this pass runs, in exactly these words: **Do not change meaning or voice anywhere in the book.** Claude flags borderline calls instead of making silent fixes. You rule each one.

Pass 5 — Read aloud. *Model: Fable for Claude's half, then you.*

Door closed, full voice, no skimming. Read the entire chapter out loud. Mark every stumble — a stumble is a sentence that trips your mouth, and it will trip your reader's eye. Fix what tripped you.

Then read the closing paragraph one more time and ask: **Is this the sentence I want my reader carrying to the next chapter?**

What you should have when you're done: A clean chapter. A banned-phrase list worth more than the chapter — because the chapter is done once and the list makes every chapter after it better.

Chapter 11 — Make It a Book: The Creative Build

What you build: The four artifacts that turn your manuscript into something a stranger can hold.

Work these in order. For anything that feeds an application or a template, download the production worksheet templates from the Marque Publishing website (or use the files in the `production_templates` folder if you have them). Give them to Claude. Mark each task complete as you go.

Build 1 — The Interior Spec and Reader Proof

The spec (one page, eight decisions):

Tell Claude: *"Research standard interior conventions for [my genre]. Produce a one-page interior spec covering trim size, typeface, margins, running heads, page numbers, chapter openings, and matter order. My north star authors are [names]."*

Fiction and nonfiction are different systems. Fiction: first-line indent, justified, no paragraph spacing. Nonfiction: block paragraphs, left-aligned, 6pt paragraph spacing. Get this wrong and a novel looks like a business report.

Decision	What to lock
Trim size	Standard for your genre (6×9 fiction, 5.5×8.5 business)
Typeface and size	A classic serif. Compare x-height, not point size. Set leading at 130-135% of type size.
Margins	Genre-standard with mirror margins for recto/verso. Account for gutter binding.
Running heads	Author/title or chapter title — pick one convention for verso and recto
Page numbers	Bottom center or outside corners
Chapter openings	Recto start (fiction standard) means blank versos between chapters — affects page count
Matter order	Front: half title, title, copyright, dedication, epigraph, TOC. Back: author's note, acknowledgments, about the author
Widows and orphans	Controlled — no single lines stranded at page tops or bottoms

Use the Interior Spec Worksheet (`02_Interior_Spec_Worksheet`) for the full template with geometry checks and trap list.

The reader proof (Build A): Your locked chapters, typeset to spec. Front and back matter as placeholders. Page count marked PROVISIONAL. Give it to at least one early reader.

Build 2 — The Cover

This is another place you can choose to partner with a designer or to do it yourself with AI. I have done both and been happy with the results in both cases. If you choose to use AI, just be prepared. It is not as simple as describing your cover and having AI build you a perfect finished cover. It will take many iterations and you should be prepared to spend a lot of time getting it right. That goes for working with a designer and doing it yourself.

Before you open any tool:

Write down what the cover must say about the book — the feeling, not the design. Set a time limit for any "just modify the existing one" approach. If it isn't easy by that limit, stop and rebuild from scratch.

Brief first. Before the designer or the AI sees anything, write the cover brief (template 09_Cover_Brief_Template): audience, tone, references, the avoid list, specs. The avoid list is the disqualifier. Have Claude draft it from Canon; you rewrite the tone and what-the-book-is sections yourself.

The sunk-cost rule: Modifying an existing cover burns more time than building one fresh. Set your exit criteria before you start.

The authenticity rule: If anything on your cover depicts something real — a place, an object, a person — verify it. Real things, verified, or nothing. The facts-file rule applies to pixels, not just words.

The wrap comes later. Your front cover becomes one panel of a three-panel wrap (front, spine, back). The spine width comes from your frozen page count — don't build the wrap until Build B is done and the count is frozen. That's Chapter 12's territory.

Build 3 — The Matter

Write all of it now. The matter is the only thing standing between you and a frozen page count. Identifiers are fixed-width — ISBNs and LCCNs don't affect page count when substituted. Only content of unknown length gates the freeze.

Component	Notes
Title page	Title, subtitle, author name, publisher/imprint
Copyright page	Use the Copyright Page Template (03_Copyright_Page_Template). One placeholder line per format's ISBN. One LCCN line. Match digit counts — a missing line is a missing line of height.
Dedication	Write it yourself. Claude's only job is to leave it alone.
Epigraph	Optional. If you have one, include it.
Table of contents	Generated from your chapter titles
Author's note	Optional. Your own words.
Acknowledgments	Real names, spelled right, of people who actually did what you're thanking them for
About the Author	Held for Build 4 (the bio test)

After writing all matter: Rebuild the interior. Freeze your page count. Write the number down. It sizes your spine.

Build 4 — The Bio Test

Open a clean Claude conversation — not your project. Ask: *"Write my author bio."* Count the fabrications. There will be fabrications.

Build your facts file — everything true about you that belongs in a bio, sourced from evidence (your resume, your LinkedIn, your publications). Save it to your project folder. Add the rule to your rules file: **Nothing about you, or as you, ships unverified.**

Write the real bio from the facts file. This bio works for the rest of your writing life — one bio, verified once, copied to every platform.

The evidence rule covers more than the bio. The author's note. The acknowledgments. The description. The bio on every platform. Anything the machine writes with your name near it. The facts file is the gate.

What you should have when you're done: An interior typeset to spec. A cover (or a cover in progress). All front and back matter written. A frozen page count. A facts file and a verified bio.

Chapter 12 — The Production Run

What you build: Your book's identity, listings, and legal existence.

Download the production worksheet templates from the Marque Publishing website (or use the files in the `production_templates` folder). Give the Production Run Sheet (`00_Production_Run_Sheet`) to Claude and tell it to run your production. The run sheet has every task across all three phases, with an owner column marking each one Claude or Author.

The pattern: Claude pre-fills the data sheets. You open the data sheet and the website side by side. You work the fields in and check each task complete.

Phase 1 — Identity and Legal

Task	Owner	What to know
Run intake	Claude	Collects title, formats, trim, channels, imprint. Does NOT collect page count — that's an output, not an input.
Reserve ISBNs from Bowker	Author	myidentifiers.com. One per format. Your own ISBNs from your own block — never a free platform ISBN. A free ISBN makes the platform your publisher of record and limits distribution.
Assign ISBNs in Bowker	Author	Irreversible. Confirm the format-to-number mapping before clicking. Log as SPENT with today's date.
Log ISBNs	Claude	Uses the ISBN Log (<code>01_ISBN_Log</code>). One log, all titles, for life. PENDING until assigned, SPENT once clicked.
Generate barcodes	Claude	Bookland EAN format. KDP supplies its own; IngramSpark and Lulu print yours.
Draft copyright page	Claude	Uses the Copyright Page Template. Every format's ISBN line present even as placeholders.
Apply for LCCN	Author	pcn.loc.gov. Print editions only. Not on the critical path — do not hold the cover behind it. Commits you to sending one copy to LOC upon publication.
Draft BISAC codes and metadata	Claude	Uses the Metadata Worksheet (<code>04_Metadata_Worksheet</code>). Max 3 codes. Primary chosen for positioning — where buyers browse, not just what the book is about.
Pre-fill copyright registration fields	Claude	Uses the Deposit and Registration sheet (<code>07_Deposit_and_Registration</code>). Assembles the values now — filing happens in Phase 3 after the book exists.

Phase 2 — Physical Build

Phases 1 and 2 run in parallel. The only hard ordering: variable-length content final → Build B → page count → spine → cover wrap.

Task	Owner	What to know
Confirm all variable-length matter final	Author	Dedication, epigraph, author's note, acknowledgments, about the author. This is the real gate.
Build B — production interior	Claude	The file that ships. Typeset from the Interior Spec Worksheet.
Interior QA	Claude	11-point checklist: ornament count, no foot ornaments, no editorial markers, recto starts, blank verso clean, running heads, non-ASCII rendering, spot-render visual check, geometry, copyright page lines, page count from PDF.
Freeze page count	Claude	Read from the built PDF, not the source document. Record as FINAL.
Compute spine width	Claude	Cream stock: pages × 0.0025 in. White stock: pages × 0.002252 in. Hardcover math differs — confirm per platform.
Substitute real identifiers	Claude	Drop real ISBNs and LCCN into copyright page. Rebuild. Confirm page count unchanged.
Build cover wraps	Claude	Uses the Cover Wrap Spec (05_Cover_Wrap_Spec). One per format per platform. Cannot build before page count is frozen.
Build EPUB	Claude	Reflowable — from normalized source, not from print PDF. EPUB copyright page omits LCCN and "Printed in" line.
Upload to platforms	Author	Uses the Platform Upload Checklist (06_Platform_Upload_Checklist). Each format's ISBN must match the log — a hardcover wrap carrying a paperback ISBN passes every dimension check but is wrong.
Set pricing	Author	Ebook royalty ceiling changes — verify the current number on the platform's live pricing page. Do not let a print price default into the ebook field.
Order and read proofs	Author	Every format. Fix what you find. Re-proof if fixes were significant.
Build author page	Author	Claim on every platform. Bio from the facts file — checked once, copied everywhere.

Phase 3 — Post-Publication Legal

Cannot start before the book physically exists. Has a hard three-month clock.

Task	Owner	What to know
Calculate deadlines	Claude	Three months from publication date. Know the date before you publish.
File copyright registration	Author	copyright.gov. Claude pre-fills every field value. You file and pay. Standard Application is the default — Single Application is illegal if anyone other than the sole author contributed. Wrong filing = fee lost, deposit consumed, effective date reset.
Print shipping slip from eCO	Author	Generated only after payment. Must physically go in the box.
Mail 2 copies to Copyright Office	Author	Best edition (hardcover if both bindings exist). Filing registration with 2 copies satisfies \$407 mandatory deposit.
Mail 1 copy to Library of Congress	Author	PCN program commitment. Separate shipment, separate address.
Update ISBN log	Claude	All ISBNs SPENT with final status.

What you should have when you're done: A book with a name, a number, and a place in the world. Listings live. Legal filings submitted or scheduled. The production machine ready to run again on the next title.

Chapter 13 — Your Machine

What you built across thirteen chapters — the complete system.

The memory layer — the part that remembers what the AI can't.

- **Canon** holds every locked decision from the thesis forward. It is the source of truth. If Canon and the prose conflict, the prose is primary.
- **The session log** holds every session's work, decisions, and reasoning, in order. Open every session by reading the last entry. Close every session by writing a new one.
- **The scar journal** holds the evidence that built your rules — what went wrong, when, and what you caught before it shipped.

The verification layer — the part that keeps the AI honest.

- **The rules file** grows with every scar. Plain-language rules Claude loads at the start of every session.
- **The preflight** loads your ground truth: rules, Canon, voice, last session log, the bar. Confirms the target. Reports back and waits for your go.
- **The postflight** asks the six questions and saves everything.
- **The verification gate** (Rule 27) runs after every draft, before you see it — a mechanical comparison of every factual claim in the prose against the preflight content. The gate catches what the instructions can't prevent, because the instruction to prefer files over recall is itself stored in recall.

The voice layer — the part that keeps the book yours.

- **Your voice profile** — drafted by Claude from your samples, rewritten by you in your words, loaded every session.
- **Your benchmark passage** — the bar, set in your own writing. Every draft measured against it.
- **Your banned-phrase list** — AI tics caught and killed, growing with every edit pass. Worth more than any single chapter.

The production layer — the part that turns a manuscript into a published book.

- **The interior spec** — one page, eight decisions, locked before typesetting.
- **The facts file** — everything true about you, with evidence. Nothing about you ships unverified.
- **The production checklist** — the owner column. Claude's lane and yours. The whole division of labor.

Your project folder — the complete inventory:

File	What it holds	Created
Canon	Every locked decision from the thesis forward	Ch 2
Session log	Every session's work, decisions, reasoning, next steps	Ch 1
Scar journal	Evidence that built your rules — failures caught	Ch 1
Rules file	Every rule (plain-language), growing with every scar. Contains preflight/postflight standing orders and the banned-phrase list	Ch 6
Voice profile	Your voice, described in your own words	Ch 7
Benchmark passage	The bar, in your own writing	Ch 8
Facts file	Everything true about you, with evidence	Ch 11
Interior spec	One page, eight decisions	Ch 11
Progress tracker	Status of every chapter	Ch 6
Production worksheets	The owner column — Claude's lane and yours	Ch 12

The machine is yours. It is standing by for the next project.

